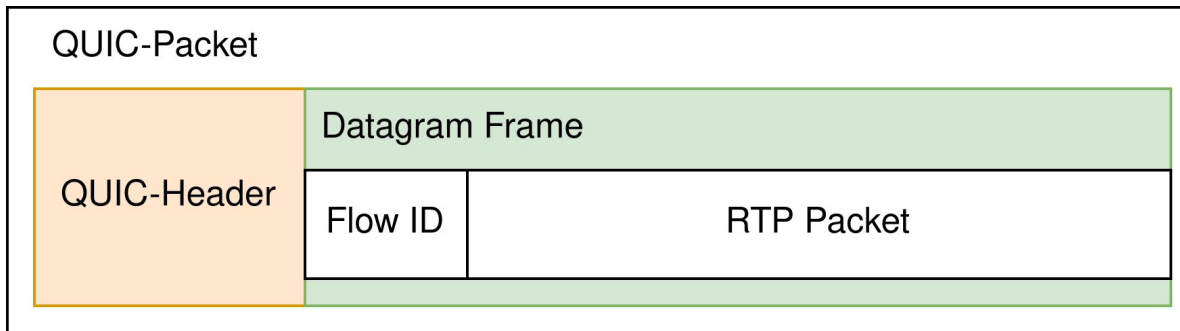


# Congestion Control for RTP over QUIC

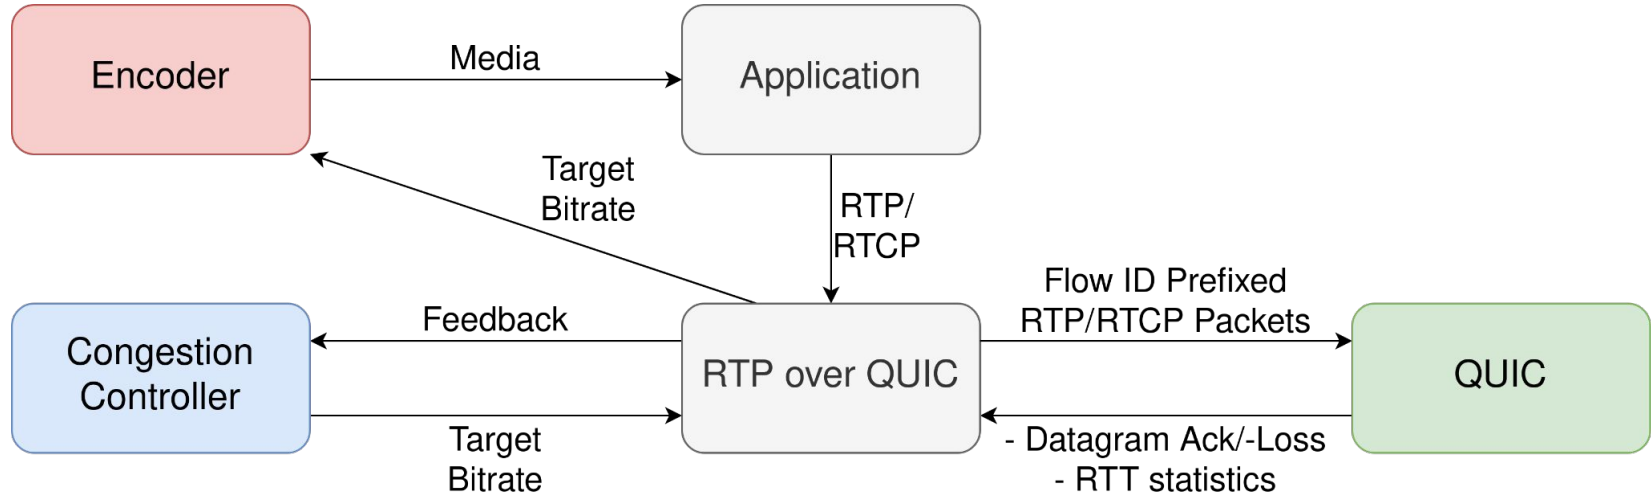
M. Engelbart and J. Ott

# RTP over QUIC

- Encapsulation for carrying RTP/RTCP over QUIC
- Using unreliable datagram extension
- Flow IDs to demultiplex datagram flows



# RTP over QUIC



- QUIC Interface Requirements (Expose Acks/Loss and RTT statistics)
- Congestion Controller Requirements  
(Feedback input / Target bandwidth output)

# Congestion Control in QUIC

- QUIC specifies congestion controller similar to TCP NewReno
- QUIC provides connection statistics to implement congestion control
  - RTT (latest, min, smoothed, mean-deviation)
  - acknowledgments
- Senders may choose a different algorithm, several under investigation
  - Cubic, BBR(v2)

# Congestion Control in RTP

- Requires low-delay while still providing useful bandwidth
- Loss-based algorithms not applicable due to filled network queues and thus large delay (-variations)
- RMCAT proposed **SCReAM** and NADA (and GCC) to use for congestion control for real-time applications
- RTP uses RTCP to signal congestion (e.g. RFC 8888)
  - Acknowledgments
  - Arrival Timestamps

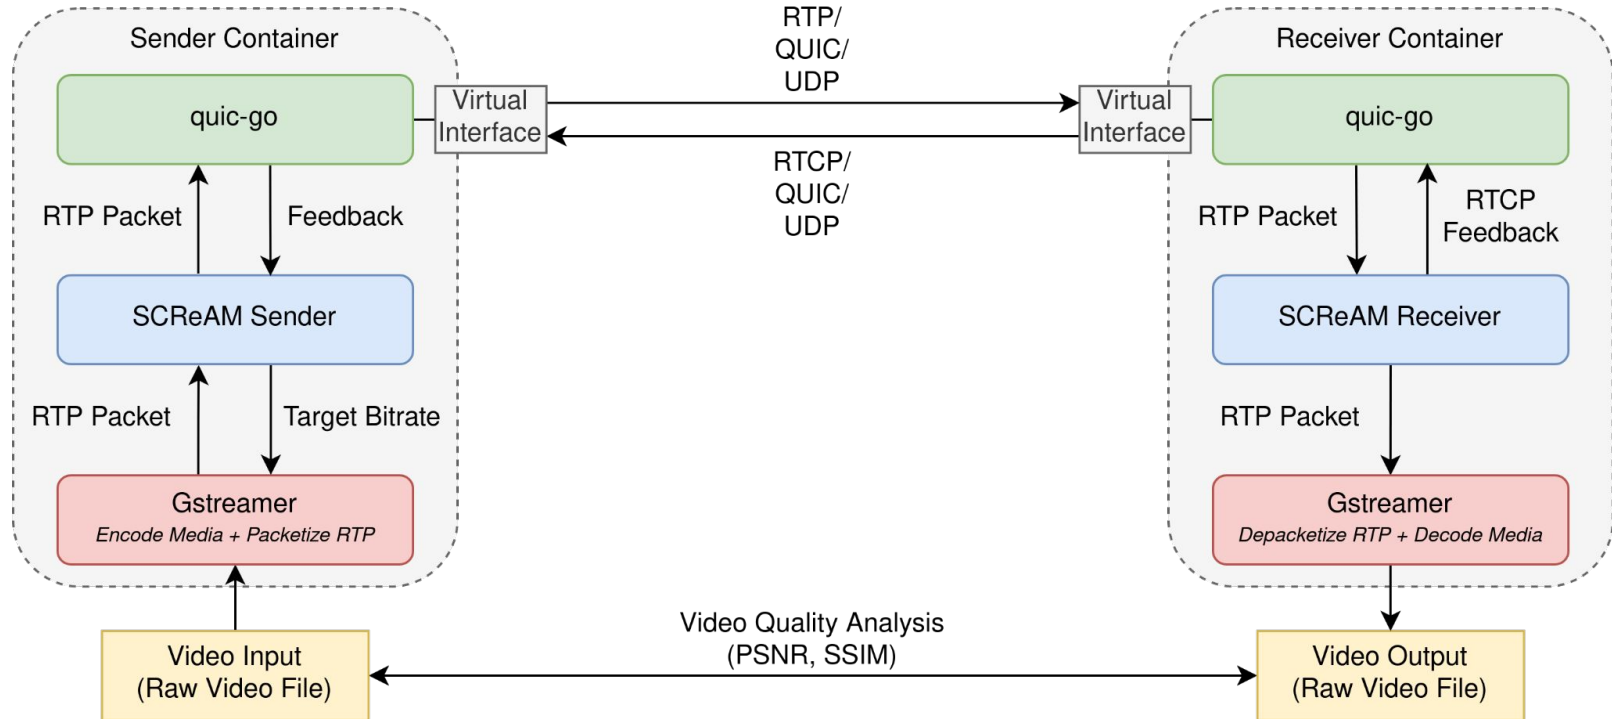
# Combining Congestion Control in RTP and QUIC

- How to do proper congestion control for interactive real-time media over QUIC given RTP realizes its own congestion control?
- Use only QUIC congestion control (NewReno) combined with trivial encoder rate control
- Disable QUIC's internal congestion control and only use real-time congestion control (SCReAM)
- Use real-time and QUIC congestion control simultaneously (SCReAM+NewReno)

# Congestion Signaling


- How to avoid duplicate signaling in RTCP?
- Instead of using RTCP, feedback can be created at the sender using QUIC connection stats
- QUIC Datagram acknowledgements signal reception of RTP packets in the ACK'ed DATAGRAM frames
- Use QUIC `latest_rtt` to infer the receive time of an RTP packet:
  - $\text{receive-ts} = \text{send-ts} + \text{latest\_rtt} / 2$

# Implementation and Testbed






# RTP over QUIC/UDP using SCReAM




SCReAM target bitrates in kbit/s




Difference of SCReAM target bitrates

# Running SCReAM over NewReno


- SCReAM target bitrate similar to RTP over UDP, when application limited
- Problematic, when both congestion controllers try to adapt to bandwidth reduction



# Only SCReAM and Reducing RTCP (50ms OWD)




RTCP Feedback




Feedback from QUIC connection statistics

# Only SReAM and Reducing RTCP (150ms OWD)



RTCP Feedback



Feedback from QUIC connection statistics

# Prioritization

- We ran all experiments again while streaming data on a QUIC stream parallel to RTP in QUIC Datagrams
- Results show, that some form of prioritization is necessary
- Without prioritization, real-time streams may degrade or even starve as a function of the internal operation of the QUIC implementation
- Test cases were rather artificial, more investigation with a more natural form of background traffic needs to be done

# Conclusion

## Main takeaways:

1. Two separate CC loops at transport and media level (expectedly) problematic
2. We can reuse QUIC state to reduce RTCP feedback
3. Prioritization is necessary

## Next steps:

- Extend testbed with new test scenarios (RFC 8867) and more CC schemes
- Come up with some prioritization scheme that gives a reasonable share of bandwidth to each real-time/non-real-time stream